

# Introduction to R Workshop (Session 1)

Dr Carlos Moreno-Garcia

School of Computing

Robert Gordon University


# What is R?

- R is an open-source programming language for statistical analysis and graphics.
- Command-line based, however, complementary tools provide a friendly user interface.
- Will require you to learn both syntax and semantic of R.




# What do we use R for?

- Exploratory and statistical data analysis.
- Visualisation and graphics.
- Data preparation (data wrangling).
- Machine learning and modelling.



# Why R?

- Free.
- Easy to use.
- Has a package for everything.
- Has a great online support community.
- Is a statistical tool AND a programming language.
- Available across platforms.
- Similar to Python and Matlab.
- Robust for visualisations.
- You can produce reports of your work easily.



Some stats


Rank	Language	Type	Score
1	Python ▼	  	100.0
2	Java ▼	  	95.3
3	C ▼	  	94.6
4	C++ ▼	  	87.0
5	JavaScript ▼		79.5
6	R ▼		78.6
7	Arduino ▼		73.2
8	Go ▼	 	73.1
9	Swift ▼	 	70.5
10	Matlab ▼		68.4

<https://spectrum.ieee.org/at-work/tech-careers/top-programming-language-2020>



## SALARY

### SALARY BY LANGUAGE




### What Languages Are Associated with the Highest Salaries Worldwide?

Global

United States

Clojure	\$90k	
F#	\$80k	
Go	\$80k	
Scala	\$78k	
Elixir	\$76k	
Ruby	\$75k	
WebAssembly	\$73k	
Rust	\$72k	
Erlang	\$71k	
Bash/Shell/PowerShell	\$69k	
R	\$64k	
Python	\$63k	
Objective-C	\$62k	
TypeScript	\$60k	
C#	\$59k	
Swift	\$59k	
Kotlin	\$57k	
SQL	\$57k	
JavaScript	\$56k	
C++	\$55k	
HTML/CSS	\$55k	

# How Technologies Are Connected





# Getting and installing R

1. Download R

<http://cran.r-project.org/>

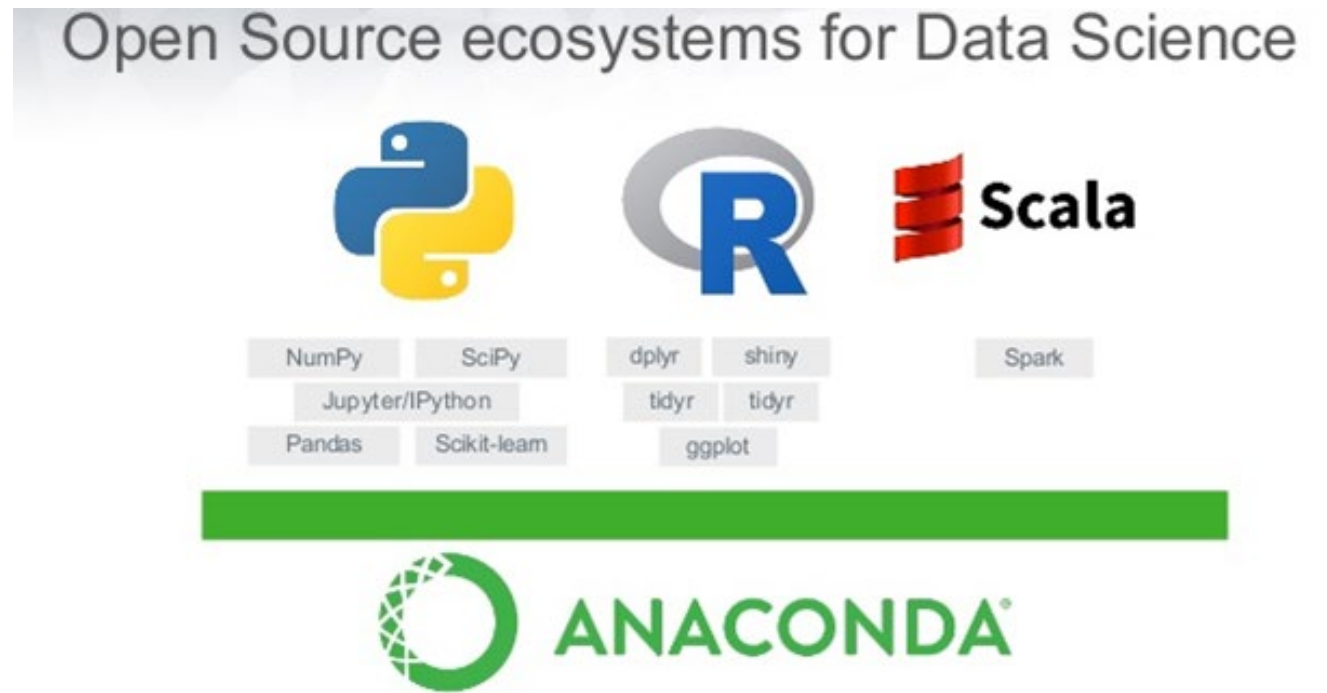
2. Download RStudio

<http://www.rstudio.com>

OR

1. Download Anaconda

<https://www.anaconda.com/>



# Introduction to R: operators & data types

- R in principle is like a calculator, you execute a command and R responds.
- These commands are mathematical operations.
- Therefore, we need to first know the operators and the data types.

# Operators

Arithmetic Operators	+	-	*	/	%%	%/%	^
Relational Operators	<	>	==	<=	>=	!=	
Logical Operators	&		!	&&			
Assignment Operators	=	<-	->	<<-	->>		
Misc. Operators	:	%in%		%*%			

<https://www.tutorialkart.com/r-tutorial/r-operators/>



# Data types

Character	String (text) values e.g. "data science"
Numeric	Decimal values e.g. 5.2
Integer	Whole numbers e.g. 7
Logical	Boolean True or False value
Factor	Categorical values e.g. employment status
Date Time	Date and time data e.g. "2015-05-12"
Complex	Complex numbers e.g. $3 + 2i$

# Variables

- A variable is a named placeholder for data.
- Initialised using the assignment operator "<-" or "=".

Character	<code>x &lt;- "data science"</code>
Numeric	<code>x &lt;- 5</code>
Logical	<code>x &lt;- T</code> Or <code>x &lt;- TRUE</code>
Factor	<code>factor("Employed", "Unemployed")</code>
Date	<code>as.Date("2015-05-12")</code>
Date Time	<code>as.POSIXct("2015-05-12 12:00")</code>



# Functions

- A function is a named group of code that is used to give instructions to R.
- May or may not accept input parameters.
- May or may not return a value.
- R comes with an extensive collection of in-built functions.
- General syntax of a function:
  - A function can be identified by the “()” after the name.
  - Example: **function.name(parameters)**.
  - A “.” in a function’s name is just there to split words!

# Examples of functions

<b>Function</b>	<b>Explanation</b>
<code>data()</code>	List all datasets currently available to R
<code>data(foo)</code>	Load dataset 'foo' into the current work space
<code>getwd()</code>	Print current working directory

# Getting help

Function	Explanation
<code>help.start()</code>	Launches the general R help in a browser which contains manuals, FAQs and reference materials
<code>help("foo")</code> or <code>?foo</code>	Help on function foo
<code>help.search("foo")</code>	Search the help system for instances of the string foo
<code>example("foo")</code>	Examples of function foo
<code>RSiteSearch("foo")</code>	Examples of the function foo in online help manuals and archived mailing list






# Packages

- Packages are named collections of (related) functions and data that are encapsulated and distributed together.

Function	Explanation
<code>install.packages('foo')</code>	Download and install package 'foo'
<code>library()</code>	List all installed packages
<code>library(foo)</code>	Load package 'foo' into the current work space

# Data structures

- R provides the following structures to collect data types/variables:



# Vectors

- Vectors are created using the `c()` function.
- general syntax is `c(val_1, val_2, ..., val_n)`.
- val can be any valid data type AS LONG AS IT REMAINS THE SAME.

Numeric	<code>x &lt;- c(1,2,3,4)</code>
Character	<code>x &lt;- c('a','b','c')</code>
Logical	<code>x &lt;- c(T,T,F,T,F)</code>



# Vectorised operations

- Operations in R are applied to entire vectors, instead of individual data elements within the vectors.

```
> x <- c(1:10)
```

```
> x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
> x * 2
```

```
[1] 2 4 6 8 10 12 14 16 18 20
```

# Data frames


- Data Frames are created using the `data.frame()` function
- general syntax is `data.frame(vect1, vect2, ..., vectn)`.

```
> names <- c('Harry', 'Bob', 'Jane')
```

```
> ages <- c(10, 9, 7)
```

```
> records <- data.frame(names, ages)
```

# Loading data into R





# Loading data from packages

- R comes pre-installed with a number of datasets.
- Third-party packages also come with more datasets.
- Use the function `data()` to see a list of installed datasets.
- Load the 'Sonar' dataset from the 'mlbench' library:

```
> library(mlbench)  
> data(Sonar)
```

- This dataset contains 208 observations of the classification of sonar signals (Mines vs Rocks).

<https://www.rdocumentation.org/packages/mlbench/versions/2.1-1/topics/Sonar>



# Loading Data from .csv files

- Delimited files can be read in R using the `read.table()` function.
- Comma separated value (CSV) files can also be read using the `read.csv()` function.

```
> dataFrame <- read.table("C:/r/data.csv", header=TRUE, sep=",")  
> dataFrame <- read.csv("C:/r/data.csv")
```



# Exploring data


Function	Description
<code>head(dataset, n )</code>	Show top n rows of dataset
<code>tail(dataset, n )</code>	Show bottom n rows of dataset
<code>ncol(dataset)</code>	Show number of columns of dataset
<code>nrow(dataset)</code>	Show number of rows of dataset
<code>dim(dataset)</code>	Show dimensions of dataset

# Exploring data

Function	Description
<code>names(dataset)</code>	List names of columns of dataset
<code>str(dataset)</code>	Show structure of dataset
<code>summary(dataset)</code>	Show summary of columns of dataset
<code>size(dataset)</code>	Show size (in bytes) of dataset

# Sub-setting data frames

- R provides a number of ways for accessing parts or elements of a data frame.
- New data frames can be generated from these subsections for:
  - Better analysis
  - Machine learning



# Four main ways of sub-setting

[ ]	Returns a subset of an object
subset()	Same as [ ] but different syntax
[[ ]]	Returns elements of a object
\$	Returns elements of a object referenced by name

# Selecting an object

- Select the first reading (i.e. row) of the Sonar dataset.

```
> sonar[1,0]  
      v1      v2      v3      v4      v5      v6      v7      v8      v9      v10     v11     v12  
1 0.02 0.0371 0.0428 0.0207 0.0954 0.0986 0.1539 0.1601 0.3109 0.2111 0.1609 0.1582  
      v13      v14      v15      v16      v17      v18      v19      v20      v21      v22      v23      v24  
1 0.2238 0.0645 0.066 0.2273 0.31 0.2999 0.5078 0.4797 0.5783 0.5071 0.4328 0.555  
      v25      v26      v27      v28      v29      v30      v31      v32      v33      v34      v35      v36  
1 0.6711 0.6415 0.7104 0.808 0.6791 0.3857 0.1307 0.2604 0.5121 0.7547 0.8537 0.8507  
      v37      v38      v39      v40      v41      v42      v43      v44      v45      v46      v47      v48  
1 0.6692 0.6097 0.4943 0.2744 0.051 0.2834 0.2825 0.4256 0.2641 0.1386 0.1051 0.1343  
      v49      v50      v51      v52      v53      v54      v55      v56      v57      v58      v59      v60  
1 0.0383 0.0324 0.0232 0.0027 0.0065 0.0159 0.0072 0.0167 0.018 0.0084 0.009 0.0032  
      class  
1      R
```



# Selecting specific rows/columns

- Select rows 10, 20, 30 and columns 1, 3, 4, 6.

```
> Sonar[c(10,20,30), c(1,3,4,6)]
```

	V1	V3	V4	V6
10	0.0164	0.0347	0.0070	0.0671
20	0.0126	0.0641	0.1732	0.2559
30	0.0189	0.0197	0.0622	0.0789



# Sub-setting by range of rows/columns

- Select rows 10-30 and columns 1-5.

```
> head(Sonar[10:20, 1:5])
```

	V1	V2	V3	V4	V5
10	0.0164	0.0173	0.0347	0.0070	0.0187
11	0.0039	0.0063	0.0152	0.0336	0.0310
12	0.0123	0.0309	0.0169	0.0313	0.0358
13	0.0079	0.0086	0.0055	0.0250	0.0344
14	0.0090	0.0062	0.0253	0.0489	0.1197
15	0.0124	0.0433	0.0604	0.0449	0.0597



# Excluding specific rows/columns

- Select rows 4 to 13 and the last 5 columns.

```
> head(Sonar[-(1:3), -(1:55)], n=10)
```

	V56	V57	V58	V59	V60	class
4	0.0073	0.0050	0.0044	0.0040	0.0117	R
5	0.0015	0.0072	0.0048	0.0107	0.0094	R
6	0.0089	0.0057	0.0027	0.0051	0.0062	R
7	0.0138	0.0092	0.0143	0.0036	0.0103	R
8	0.0097	0.0085	0.0047	0.0048	0.0053	R
9	0.0049	0.0065	0.0093	0.0059	0.0022	R
10	0.0068	0.0032	0.0035	0.0056	0.0040	R
11	0.0093	0.0042	0.0003	0.0053	0.0036	R
12	0.0118	0.0026	0.0092	0.0009	0.0044	R
13	0.0019	0.0059	0.0058	0.0059	0.0032	R



# Sub-setting by column name

- Select “the first” rows and the column labelled “V23”.

```
> head(Sonar["V23"])  
      V23  
1 0.4328  
2 0.3957  
3 0.4293  
4 0.5556  
5 0.5730  
6 0.5890
```

# Sub-setting by column name

- Select “the first” rows and the columns labelled “V23”, “V24” and “V25”.

```
> head(Sonar[names(Sonar) %in% c("V23", "V24", "V25")])
```

	V23	V24	V25
1	0.4328	0.5550	0.6711
2	0.3957	0.3914	0.3250
3	0.4293	0.3648	0.5331
4	0.5556	0.4846	0.3140
5	0.5730	0.5399	0.3161
6	0.5890	0.2872	0.2043

- You can exclude columns by using !names() instead.



# Accessing the content of a column

- Select “the first” columns labelled “V23”.

```
> head(Sonar[["v23"]])  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890  
> head(Sonar[[23]])  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890  
> head(Sonar$v23)  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890
```

- What is the difference between [], \$ and [[]]?

```
> class(Sonar["v23"])  
[1] "data.frame"  
> class(Sonar[[23]])  
[1] "numeric"  
> class(Sonar[["v23"]])  
[1] "numeric"  
> class(Sonar$"v23")  
[1] "numeric"
```

# Select rows that meet a condition

- Select rows where "V23" is greater than 0.96.

```
> Sonar$V23>0.96
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[14] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[27] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE TRUE TRUE FALSE FALSE
[40] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[53] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
[66] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE
[79] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[92] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[105] FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[118] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE
[131] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE
[144] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[170] FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[183] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[196] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
> Sonar[Sonar$V23>0.96,22:25]
```

	V22	V23	V24	V25
34	0.8982	0.9664	0.8515	0.6626
36	0.4876	1.0000	0.8675	0.4718
37	0.8793	0.9606	0.8786	0.6905
62	0.8747	1.0000	0.8948	0.8420
77	0.9976	0.9872	0.9761	0.9009
78	0.9814	0.9620	0.9601	0.9118
80	1.0000	0.9645	0.9432	0.8658
110	0.9422	1.0000	0.9931	0.9575
111	0.9338	1.0000	0.9102	0.8496
122	0.9668	1.0000	0.9893	0.9376
126	0.8537	0.9642	1.0000	0.9357
128	0.9473	1.0000	0.8975	0.7806
138	0.8454	0.9739	1.0000	0.6665
139	0.6572	0.9734	0.9757	0.8079
175	0.9385	1.0000	0.9831	0.9932

# Select rows that meet two conditions

- Select rows where "V23" is greater than 0.96 AND V24 equal to 1.

```
> Sonar[Sonar$V23>0.96 & Sonar$V24==1,22:25]
```

	V22	V23	V24	V25
126	0.8537	0.9642	1	0.9357
138	0.8454	0.9739	1	0.6665

# Select rows that meet two conditions

- Select rows where "V23" is greater than 0.96 OR V24 equal to 1.

```
> Sonar[Sonar$V23>0.96 | Sonar$V24==1,22:25]
      V22      V23      V24      V25
34  0.8982 0.9664 0.8515 0.6626
36  0.4876 1.0000 0.8675 0.4718
37  0.8793 0.9606 0.8786 0.6905
62  0.8747 1.0000 0.8948 0.8420
76  0.9403 0.9409 1.0000 0.9725
77  0.9976 0.9872 0.9761 0.9009
78  0.9814 0.9620 0.9601 0.9118
80  1.0000 0.9645 0.9432 0.8658
86  0.7545 0.8311 1.0000 0.8762
87  0.7569 0.8596 1.0000 0.8457
90  0.6794 0.8297 1.0000 0.8240
110 0.9422 1.0000 0.9931 0.9575
111 0.9338 1.0000 0.9102 0.8496
122 0.9668 1.0000 0.9893 0.9376
126 0.8537 0.9642 1.0000 0.9357
128 0.9473 1.0000 0.8975 0.7806
138 0.8454 0.9739 1.0000 0.6665
139 0.6572 0.9734 0.9757 0.8079
175 0.9385 1.0000 0.9831 0.9932
201 0.7924 0.8793 1.0000 0.9865
```



# Lab Activity Session 1

- Open RStudio and load “GSM0008\_S1\_Lab.R”
- Follow the script to practice what we have learnt so far.