

Introduction to R Workshop (Session 1)

Dr Carlos Moreno-Garcia
School of Computing
Robert Gordon University


What is R?

- R is an open-source programming language for statistical analysis and graphics.
- Command-line based, however, complementary tools provide a friendly user interface.
- Will require you to learn both syntax and semantic of R.




What do we use R for?

- Exploratory and statistical data analysis.
- Visualisation and graphics.
- Data preparation (data wrangling).
- Machine learning and modelling.















Why R?

- Free.
- Easy to use.
- Has a package for everything.
- Has a great online support community.
- Is a statistical tool AND a programming language.
- Available across platforms.
- Similar to Python and MATLAB.
- Robust for visualisations.
- You can produce reports of your work easily.




Some stats

Rank	Language	Type	Score
1	Python	  	100.0
2	Java	  	95.4
3	C	  	94.7
4	C++	  	92.4
5	JavaScript		88.1
6	C#	   	82.4
7	R		81.7
8	Go	 	77.7
9	HTML		75.4
10	Swift	 	70.4

<https://spectrum.ieee.org/top-programming-languages/>

Top paying technologies




<https://insights.stackoverflow.com/survey/2021>


Salary and experience by language



PHP developers are disproportionately underpaid compared to other languages with the same experience.



How Technologies Are Connected



Getting and installing R

1. Download R

<http://cran.r-project.org/>

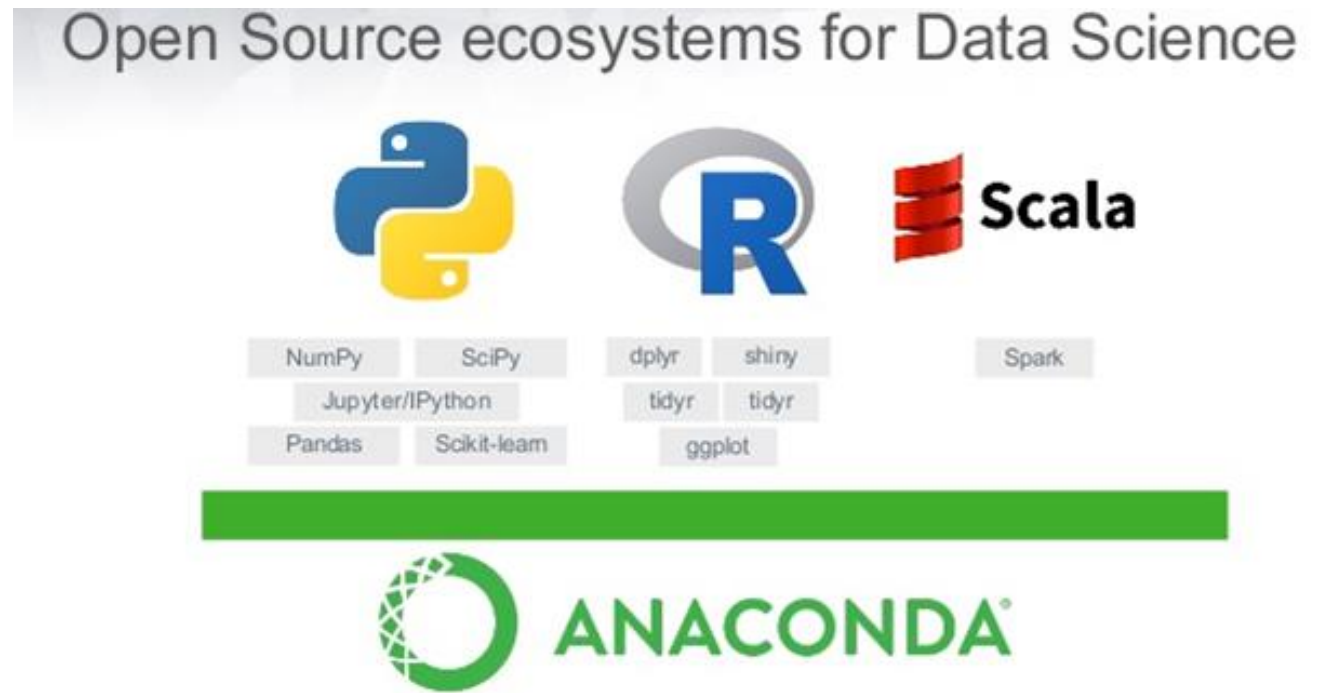
2. Download RStudio

<http://www.rstudio.com>

OR

1. Download Anaconda

<https://www.anaconda.com/>



Introduction to R: operators & data types

- R in principle is like a calculator, you execute a command and R responds.
- These commands are mathematical operations.
- Therefore, we need to first know the operators and the data types.

Operators

Arithmetic Operators	+	-	*	/	%%	%/%%	^
Relational Operators	<	>	==	<=	>=	!=	
Logical Operators	&		!	&&			
Assignment Operators	=	<-	->	<<-	->>		
Misc. Operators	:	%in%		%*%			

<https://www.tutorialkart.com/r-tutorial/r-operators/>

Data types

Character	String (text) values e.g. "data science"
Numeric	Decimal values e.g. 5.2
Integer	Whole numbers e.g. 7
Logical	Boolean True or False value
Factor	Categorical values e.g. employment status
Date Time	Date and time data e.g. "2015-05-12"
Complex	Complex numbers e.g. $3 + 2i$

Variables

- A variable is a named placeholder for data.
- Initialised using the assignment operator "<-" or "=".

Character	<code>x <- "data science"</code>
Numeric	<code>x <- 5</code>
Logical	<code>x <- T</code> Or <code>x <- TRUE</code>
Factor	<code>factor("Employed", "Unemployed")</code>
Date	<code>as.Date("2015-05-12")</code>
Date Time	<code>as.POSIXct("2015-05-12 12:00")</code>

Functions

- A function is a named group of code that is used to give instructions to R.
- May or may not accept input parameters.
- May or may not return a value.
- R comes with an extensive collection of in-built functions.
- General syntax of a function:
 - A function can be identified by the “()” after the name.
 - Example: **function.name(parameters)**.
 - A “.” in a function’s name is just there to split words!

Examples of functions

Function	Explanation
<code>data()</code>	List all datasets currently available to R
<code>data(foo)</code>	Load dataset 'foo' into the current work space
<code>getwd()</code>	Print current working directory

Getting help

Function	Explanation
<code>help.start()</code>	Launches the general R help in a browser which contains manuals, FAQs and reference materials
<code>help("foo")</code> or <code>?foo</code>	Help on function foo
<code>help.search("foo")</code>	Search the help system for instances of the string foo
<code>example("foo")</code>	Examples of function foo
<code>RSiteSearch("foo")</code>	Examples of the function foo in online help manuals and archived mailing list


Packages

- Packages are named collections of (related) functions and data that are encapsulated and distributed together.

Function	Explanation
<code>install.packages('foo')</code>	Download and install package 'foo'
<code>library()</code>	List all installed packages
<code>library(foo)</code>	Load package 'foo' into the current work space

Data structures

- R provides the following structures to collect data types/variables:



Vectors

- Vectors are created using the `c()` function.
- general syntax is `c(val_1, val_2, ..., val_n)`.
- *val* can be any valid data type AS LONG AS IT REMAINS THE SAME.

Numeric	<code>x <- c(1,2,3,4)</code>
Character	<code>x <- c('a','b','c')</code>
Logical	<code>x <- c(T,T,F,T,F)</code>

Vectorised operations

- Operations in R are applied to entire vectors, instead of individual data elements within the vectors.

```
> x <- c(1:10)
```

```
> x
```

```
[1]  1  2  3  4  5  6  7  8  9 10
```

```
> x * 2
```

```
[1]  2  4  6  8 10 12 14 16 18 20
```

Data frames


- Data Frames are created using the `data.frame()` function
- general syntax is `data.frame(vect1, vect2, ..., vectn)`.

```
> names <- c('Harry', 'Bob', 'Jane')
```

```
> ages <- c(10, 9, 7)
```

```
> records <- data.frame(names, ages)
```

Loading data into R



Loading data from packages

- R comes pre-installed with a number of datasets.
- Third-party packages also come with more datasets.
- Use the function `data()` to see a list of installed datasets.
- Load the 'Sonar' dataset from the 'mlbench' library:

```
> library(mlbench)  
> data(Sonar)
```

- This dataset contains 208 observations of the classification of sonar signals (Mines vs Rocks).

<https://www.rdocumentation.org/packages/mlbench/versions/2.1-1/topics/Sonar>

Loading Data from .csv files

- Delimited files can be read in R using the `read.table()` function.
- Comma separated value (CSV) files can also be read using the `read.csv()` function.

```
> dataFrame <- read.table("C:/r/data.csv", header=TRUE, sep=",")  
> dataFrame <- read.csv("C:/r/data.csv")
```

Exploring data


Function	Description
<code>head(dataset, n)</code>	Show top n rows of dataset
<code>tail(dataset, n)</code>	Show bottom n rows of dataset
<code>ncol(dataset)</code>	Show number of columns of dataset
<code>nrow(dataset)</code>	Show number of rows of dataset
<code>dim(dataset)</code>	Show dimensions of dataset

Exploring data

Function	Description
<code>names(dataset)</code>	List names of columns of dataset
<code>str(dataset)</code>	Show structure of dataset
<code>summary(dataset)</code>	Show summary of columns of dataset
<code>size(dataset)</code>	Show size (in bytes) of dataset

Sub-setting data frames

- R provides a number of ways for accessing parts or elements of a data frame.
- New data frames can be generated from these subsections for:
 - Better analysis
 - Machine learning



Four main ways of sub-setting

[]	Returns a subset of an object
subset()	Same as [] but different syntax
[[]]	Returns elements of a object
\$	Returns elements of a object referenced by name

Selecting an object

- Select the first reading (i.e. row) of the Sonar dataset.

```
> sonar[1,]  
      v1      v2      v3      v4      v5      v6      v7      v8      v9      v10     v11     v12  
1 0.02 0.0371 0.0428 0.0207 0.0954 0.0986 0.1539 0.1601 0.3109 0.2111 0.1609 0.1582  
      v13     v14     v15     v16     v17     v18     v19     v20     v21     v22     v23     v24  
1 0.2238 0.0645 0.066 0.2273 0.31 0.2999 0.5078 0.4797 0.5783 0.5071 0.4328 0.555  
      v25     v26     v27     v28     v29     v30     v31     v32     v33     v34     v35     v36  
1 0.6711 0.6415 0.7104 0.808 0.6791 0.3857 0.1307 0.2604 0.5121 0.7547 0.8537 0.8507  
      v37     v38     v39     v40     v41     v42     v43     v44     v45     v46     v47     v48  
1 0.6692 0.6097 0.4943 0.2744 0.051 0.2834 0.2825 0.4256 0.2641 0.1386 0.1051 0.1343  
      v49     v50     v51     v52     v53     v54     v55     v56     v57     v58     v59     v60  
1 0.0383 0.0324 0.0232 0.0027 0.0065 0.0159 0.0072 0.0167 0.018 0.0084 0.009 0.0032  
      class  
1      R
```

Selecting specific rows/columns

- Select rows 10, 20, 30 and columns 1, 3, 4, 6.

```
> Sonar[c(10,20,30), c(1,3,4,6)]
```

	V1	V3	V4	V6
10	0.0164	0.0347	0.0070	0.0671
20	0.0126	0.0641	0.1732	0.2559
30	0.0189	0.0197	0.0622	0.0789

Sub-setting by range of rows/columns

- Select rows 10-30 and columns 1-5.

```
> head(Sonar[10:20, 1:5])
```

	V1	V2	V3	V4	V5
10	0.0164	0.0173	0.0347	0.0070	0.0187
11	0.0039	0.0063	0.0152	0.0336	0.0310
12	0.0123	0.0309	0.0169	0.0313	0.0358
13	0.0079	0.0086	0.0055	0.0250	0.0344
14	0.0090	0.0062	0.0253	0.0489	0.1197
15	0.0124	0.0433	0.0604	0.0449	0.0597

Excluding specific rows/columns

- Select rows 4 to 13 and the last 5 columns.

```
> head(Sonar[-(1:3), -(1:55)], n=10)
```

	V56	V57	V58	V59	V60	class
4	0.0073	0.0050	0.0044	0.0040	0.0117	R
5	0.0015	0.0072	0.0048	0.0107	0.0094	R
6	0.0089	0.0057	0.0027	0.0051	0.0062	R
7	0.0138	0.0092	0.0143	0.0036	0.0103	R
8	0.0097	0.0085	0.0047	0.0048	0.0053	R
9	0.0049	0.0065	0.0093	0.0059	0.0022	R
10	0.0068	0.0032	0.0035	0.0056	0.0040	R
11	0.0093	0.0042	0.0003	0.0053	0.0036	R
12	0.0118	0.0026	0.0092	0.0009	0.0044	R
13	0.0019	0.0059	0.0058	0.0059	0.0032	R

Sub-setting by column name

- Select “the first” rows and the column labelled “V23”.

```
> head(Sonar["V23"])  
      V23  
1 0.4328  
2 0.3957  
3 0.4293  
4 0.5556  
5 0.5730  
6 0.5890
```

Sub-setting by column name

- Select “the first” rows and the columns labelled “V23”, “V24” and “V25”.

```
> head(Sonar[names(Sonar) %in% c("V23", "V24", "V25")])
```

	V23	V24	V25
1	0.4328	0.5550	0.6711
2	0.3957	0.3914	0.3250
3	0.4293	0.3648	0.5331
4	0.5556	0.4846	0.3140
5	0.5730	0.5399	0.3161
6	0.5890	0.2872	0.2043

- You can exclude columns by using !names() instead.

Accessing the content of a column

- Select “the first” columns labelled “V23”.

```
> head(Sonar[["v23"]])  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890  
> head(Sonar[[23]])  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890  
> head(Sonar$v23)  
[1] 0.4328 0.3957 0.4293 0.5556 0.5730 0.5890
```

- What is the difference between [], \$ and [[]]?

```
> class(Sonar["v23"])  
[1] "data.frame"  
> class(Sonar[[23]])  
[1] "numeric"  
> class(Sonar[["v23"]])  
[1] "numeric"  
> class(Sonar$"v23")  
[1] "numeric"
```

Select rows that meet a condition

- Select rows where "V23" is greater than 0.96.

```
> Sonar$V23>0.96
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[14] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[27] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE TRUE TRUE FALSE FALSE
[40] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[53] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
[66] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE
[79] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[92] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[105] FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[118] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE
[131] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE
[144] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[170] FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[183] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[196] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
> Sonar[Sonar$V23>0.96,22:25]
```

	V22	V23	V24	V25
34	0.8982	0.9664	0.8515	0.6626
36	0.4876	1.0000	0.8675	0.4718
37	0.8793	0.9606	0.8786	0.6905
62	0.8747	1.0000	0.8948	0.8420
77	0.9976	0.9872	0.9761	0.9009
78	0.9814	0.9620	0.9601	0.9118
80	1.0000	0.9645	0.9432	0.8658
110	0.9422	1.0000	0.9931	0.9575
111	0.9338	1.0000	0.9102	0.8496
122	0.9668	1.0000	0.9893	0.9376
126	0.8537	0.9642	1.0000	0.9357
128	0.9473	1.0000	0.8975	0.7806
138	0.8454	0.9739	1.0000	0.6665
139	0.6572	0.9734	0.9757	0.8079
175	0.9385	1.0000	0.9831	0.9932

Select rows that meet two conditions

- Select rows where "V23" is greater than 0.96 AND V24 equal to 1.

```
> Sonar[Sonar$V23>0.96 & Sonar$V24==1,22:25]
```

	V22	V23	V24	V25
126	0.8537	0.9642	1	0.9357
138	0.8454	0.9739	1	0.6665

Select rows that meet two conditions

- Select rows where "V23" is greater than 0.96 OR V24 equal to 1.

```
> Sonar[Sonar$V23>0.96 | Sonar$V24==1,22:25]
      V22      V23      V24      V25
34  0.8982 0.9664 0.8515 0.6626
36  0.4876 1.0000 0.8675 0.4718
37  0.8793 0.9606 0.8786 0.6905
62  0.8747 1.0000 0.8948 0.8420
76  0.9403 0.9409 1.0000 0.9725
77  0.9976 0.9872 0.9761 0.9009
78  0.9814 0.9620 0.9601 0.9118
80  1.0000 0.9645 0.9432 0.8658
86  0.7545 0.8311 1.0000 0.8762
87  0.7569 0.8596 1.0000 0.8457
90  0.6794 0.8297 1.0000 0.8240
110 0.9422 1.0000 0.9931 0.9575
111 0.9338 1.0000 0.9102 0.8496
122 0.9668 1.0000 0.9893 0.9376
126 0.8537 0.9642 1.0000 0.9357
128 0.9473 1.0000 0.8975 0.7806
138 0.8454 0.9739 1.0000 0.6665
139 0.6572 0.9734 0.9757 0.8079
175 0.9385 1.0000 0.9831 0.9932
201 0.7924 0.8793 1.0000 0.9865
```

Lab Activity Session 1

- Install Rstudio!
- Open RStudio and load “GSM0008_S1_Lab.R”
- Follow the script to practice what we have learnt so far.